

Introduction

Francisco Villamil

Applied Quantitative Methods II
MA in Social Sciences, Spring 2026

Course overview

- This is the second part of the quantitative methods sequence

Course overview

- This is the second part of the quantitative methods sequence
- Focus on **applying** statistical tools in practice

Course overview

- This is the second part of the quantitative methods sequence
- Focus on **applying** statistical tools in practice
- Less theory, more hands-on work with data

Course overview

- This is the second part of the quantitative methods sequence
- Focus on **applying** statistical tools in practice
- Less theory, more hands-on work with data
- Goal: go from research question to answer

What will you learn?

- How to choose the right model for your question

What will you learn?

- How to choose the right model for your question
- How to interpret and visualize model results

What will you learn?

- How to choose the right model for your question
- How to interpret and visualize model results
- How to evaluate whether a model is appropriate

What will you learn?

- How to choose the right model for your question
- How to interpret and visualize model results
- How to evaluate whether a model is appropriate
- How to work with different types of data (panel, spatial, etc.)

What will you learn?

- How to choose the right model for your question
- How to interpret and visualize model results
- How to evaluate whether a model is appropriate
- How to work with different types of data (panel, spatial, etc.)
- Best practices in computing and reproducibility

Course structure

Feb 5	Introduction
Feb 12	Applied regression
Feb 19	Applied regression II (binary)
Feb 26	Interpretation and diagnostics
Mar 5	Best practices in computing
Mar 12	Panel data I
Mar 19	Panel data II
Mar 26	Spatial data
<i>Easter break</i>	
Apr 9	Spatial data
Apr 16	Other outcomes
Apr 23	Project presentations
Apr 30	Exam + Review

Course structure

Feb 5	Introduction
Feb 12	Applied regression
Feb 19	Applied regression II (binary)
Feb 26	Interpretation and diagnostics
Mar 5	Best practices in computing
Mar 12	Panel data I
Mar 19	Panel data II
Mar 26	Spatial data
<i>Easter break</i>	
Apr 9	Spatial data
Apr 16	Other outcomes
Apr 23	Project presentations
Apr 30	Exam + Review

Course structure

Feb 5	Introduction
Feb 12	Applied regression
Feb 19	Applied regression II (binary)
Feb 26	Interpretation and diagnostics
Mar 5	Best practices in computing
Mar 12	Panel data I
Mar 19	Panel data II
Mar 26	Spatial data
<i>Easter break</i>	
Apr 9	Spatial data
Apr 16	Other outcomes
Apr 23	Project presentations
Apr 30	Exam + Review

Course structure

Feb 5	Introduction
Feb 12	Applied regression
Feb 19	Applied regression II (binary)
Feb 26	Interpretation and diagnostics
Mar 5	Best practices in computing
Mar 12	Panel data I
Mar 19	Panel data II
Mar 26	Spatial data
<i>Easter break</i>	
Apr 9	Spatial data
Apr 16	Other outcomes
Apr 23	Project presentations
Apr 30	Exam + Review

Course structure

Feb 5	Introduction
Feb 12	Applied regression
Feb 19	Applied regression II (binary)
Feb 26	Interpretation and diagnostics
Mar 5	Best practices in computing <i>(move just before break?)</i>
Mar 12	Panel data I
Mar 19	Panel data II
Mar 26	Spatial data
<i>Easter break</i>	
Apr 9	Spatial data
Apr 16	Other outcomes
Apr 23	Project presentations
Apr 30	Exam + Review

Course structure

Feb 5	Introduction
Feb 12	Applied regression
Feb 19	Applied regression II (binary)
Feb 26	Interpretation and diagnostics
Mar 5	Best practices in computing
Mar 12	Panel data I
Mar 19	Panel data II
Mar 26	Spatial data
<i>Easter break</i>	
Apr 9	Spatial data
Apr 16	Other outcomes
Apr 23	Project presentations
Apr 30	Exam + Review

Course structure

Feb 5	Introduction
Feb 12	Applied regression
Feb 19	Applied regression II (binary)
Feb 26	Interpretation and diagnostics
Mar 5	Best practices in computing
Mar 12	Panel data I
Mar 19	Panel data II
Mar 26	Spatial data
<i>Easter break</i>	
Apr 9	Spatial data
Apr 16	Other outcomes
Apr 23	Project presentations
Apr 30	Exam + Review

Course structure

Feb 5	Introduction
Feb 12	Applied regression
Feb 19	Applied regression II (binary)
Feb 26	Interpretation and diagnostics
Mar 5	Best practices in computing
Mar 12	Panel data I
Mar 19	Panel data II
Mar 26	Spatial data
<i>Easter break</i>	
Apr 9	Spatial data
Apr 16	Other outcomes
Apr 23	Project presentations
Apr 30	Exam + Review

Course structure

Feb 5	Introduction
Feb 12	Applied regression
Feb 19	Applied regression II (binary)
Feb 26	Interpretation and diagnostics
Mar 5	Best practices in computing
Mar 12	Panel data I
Mar 19	Panel data II
Mar 26	Spatial data
<i>Easter break</i>	
Apr 9	Spatial data
Apr 16	Other outcomes
Apr 23	Project presentations
Apr 30	Exam + Review

Course structure

Feb 5	Introduction
Feb 12	Applied regression
Feb 19	Applied regression II (binary)
Feb 26	Interpretation and diagnostics
Mar 5	Best practices in computing
Mar 12	Panel data I
Mar 19	Panel data II
Mar 26	Spatial data
<i>Easter break</i>	
Apr 9	Spatial data
Apr 16	Other outcomes
Apr 23	Project presentations
Apr 30	Exam + Review

Evaluation

- Problem sets (20%)
 - Started in class, finished at home
 - Short deadlines
- Proposal presentation and peer review (10% + 10%)
- Final essay (30%)
 - Small research note (max 3,000 words)
 - Original data analysis using R
- Exam (30%)

Roadmap

The Big Picture

Workflow Basics

The research process

Theory $\longleftrightarrow$ Data Generating Process $\longleftrightarrow$ Data

- Theories make claims about how the world works

The research process

Theory $\longleftrightarrow$ Data Generating Process $\longleftrightarrow$ Data

- Theories make claims about how the world works
- These claims imply certain patterns in data

The research process

Theory $\longleftrightarrow$ Data Generating Process $\longleftrightarrow$ Data

- Theories make claims about how the world works
- These claims imply certain patterns in data
- We observe data and try to learn about the underlying process

The research process

Theory $\longleftrightarrow$ Data Generating Process $\longleftrightarrow$ Data

- Theories make claims about how the world works
- These claims imply certain patterns in data
- We observe data and try to learn about the underlying process
- Our research strategy connects theory to data

Theory first, methods second

- The research question and theory should drive everything:

Theory first, methods second

- The research question and theory should drive everything:
 - What unit of analysis to use

Theory first, methods second

- The research question and theory should drive everything:
 - What unit of analysis to use
 - What variation to look at

Theory first, methods second

- The research question and theory should drive everything:
 - What unit of analysis to use
 - What variation to look at
 - What empirical strategy to follow

Theory first, methods second

- The research question and theory should drive everything:
 - What unit of analysis to use
 - What variation to look at
 - What empirical strategy to follow
- Methods are tools to implement that strategy

Theory first, methods second

- The research question and theory should drive everything:
 - What unit of analysis to use
 - What variation to look at
 - What empirical strategy to follow
- Methods are tools to implement that strategy
- Common mistake: picking a method and then looking for a question

Theory first, methods second

- The research question and theory should drive everything:
 - What unit of analysis to use
 - What variation to look at
 - What empirical strategy to follow
- Methods are tools to implement that strategy
- Common mistake: picking a method and then looking for a question
- In this course: we learn tools, but always ask *why this tool for this question?*

What is a Data Generating Process (DGP)?

- The rules that govern how data comes to exist

What is a Data Generating Process (DGP)?

- The rules that govern how data comes to exist
- Includes:

What is a Data Generating Process (DGP)?

- The rules that govern how data comes to exist
- Includes:
 - The social or political process we study

What is a Data Generating Process (DGP)?

- The rules that govern how data comes to exist
- Includes:
 - The social or political process we study
 - How observations end up in our dataset

What is a Data Generating Process (DGP)?

- The rules that govern how data comes to exist
- Includes:
 - The social or political process we study
 - How observations end up in our dataset
- We never observe the DGP directly

What is a Data Generating Process (DGP)?

- The rules that govern how data comes to exist
- Includes:
 - The social or political process we study
 - How observations end up in our dataset
- We never observe the DGP directly
- We use statistical models to make inferences about it

Why do we need statistics?

- Our theories deal with processes, not just data

Why do we need statistics?

- Our theories deal with processes, not just data
- Data is a window into the underlying process

Why do we need statistics?

- Our theories deal with processes, not just data
- Data is a window into the underlying process
- Statistics helps us:

Why do we need statistics?

- Our theories deal with processes, not just data
- Data is a window into the underlying process
- Statistics helps us:
 - Separate signal from noise

Why do we need statistics?

- Our theories deal with processes, not just data
- Data is a window into the underlying process
- Statistics helps us:
 - Separate signal from noise
 - Quantify uncertainty

Why do we need statistics?

- Our theories deal with processes, not just data
- Data is a window into the underlying process
- Statistics helps us:
 - Separate signal from noise
 - Quantify uncertainty
 - Make valid inferences

Sources of uncertainty

- **Sampling uncertainty:** We observe a sample, not the population

Sources of uncertainty

- **Sampling uncertainty:** We observe a sample, not the population
- **Theoretical uncertainty:** Our theories are simplifications

Sources of uncertainty

- **Sampling uncertainty:** We observe a sample, not the population
- **Theoretical uncertainty:** Our theories are simplifications
- **Fundamental uncertainty:** Some processes are inherently random

Sources of uncertainty

- **Sampling uncertainty:** We observe a sample, not the population
- **Theoretical uncertainty:** Our theories are simplifications
- **Fundamental uncertainty:** Some processes are inherently random

Sources of uncertainty

- **Sampling uncertainty:** We observe a sample, not the population
- **Theoretical uncertainty:** Our theories are simplifications
- **Fundamental uncertainty:** Some processes are inherently random
- All of these create “noise” in our data

Sources of uncertainty

- **Sampling uncertainty:** We observe a sample, not the population
- **Theoretical uncertainty:** Our theories are simplifications
- **Fundamental uncertainty:** Some processes are inherently random

- All of these create “noise” in our data
- Statistical models help us deal with this noise

The logic of statistical inference

- **Probability theory:** Given a known process, what data will we see?

The logic of statistical inference

- **Probability theory:** Given a known process, what data will we see?

The logic of statistical inference

- **Probability theory:** Given a known process, what data will we see?
- **Statistical inference:** Given observed data, what can we learn about the process?

The logic of statistical inference

- **Probability theory:** Given a known process, what data will we see?
- **Statistical inference:** Given observed data, what can we learn about the process?

The logic of statistical inference

- **Probability theory:** Given a known process, what data will we see?
- **Statistical inference:** Given observed data, what can we learn about the process?
- We're doing the reverse: from data back to process

Roadmap

The Big Picture

Workflow Basics

Learning to use computers as tools

- World of quantitative methods is changing fast

Learning to use computers as tools

- World of quantitative methods is changing fast
 - e.g. Claude Code

Learning to use computers as tools

- World of quantitative methods is changing fast
 - e.g. Claude Code
- I think it'll be more important to be really literate with computers

Learning to use computers as tools

- World of quantitative methods is changing fast
 - e.g. Claude Code
- I think it'll be more important to be really literate with computers
- Part of this course will also involve learning how to properly use computers

Learning to use computers as tools

- World of quantitative methods is changing fast
 - e.g. Claude Code
- I think it'll be more important to be really literate with computers
- Part of this course will also involve learning how to properly use computers
 - Not using only RStudio, R Markdown, etc, but being ready to do big data-based projects

Learning to use computers as tools

- World of quantitative methods is changing fast
 - e.g. Claude Code
- I think it'll be more important to be really literate with computers
- Part of this course will also involve learning how to properly use computers
 - Not using only RStudio, R Markdown, etc, but being ready to do big data-based projects
- We'll have a session on computing, project management, etc – but today, some notes on version control

The problem: managing files over time

- Have you ever had files like this?

The problem: managing files over time

- Have you ever had files like this?

→ `thesis_v1.docx`

The problem: managing files over time

- Have you ever had files like this?
 - `thesis_v1.docx`
 - `thesis_v2_final.docx`

The problem: managing files over time

- Have you ever had files like this?
 - `thesis_v1.docx`
 - `thesis_v2_final.docx`
 - `thesis_v2_final_REAL.docx`

The problem: managing files over time

- Have you ever had files like this?
 - `thesis_v1.docx`
 - `thesis_v2_final.docx`
 - `thesis_v2_final_REAL.docx`
 - `thesis_v2_final_REAL_submitted.docx`

The problem: managing files over time

- Have you ever had files like this?
 - `thesis_v1.docx`
 - `thesis_v2_final.docx`
 - `thesis_v2_final_REAL.docx`
 - `thesis_v2_final_REAL_submitted.docx`

The problem: managing files over time

- Have you ever had files like this?
 - `thesis_v1.docx`
 - `thesis_v2_final.docx`
 - `thesis_v2_final_REAL.docx`
 - `thesis_v2_final_REAL_submitted.docx`
- What changed between versions?

The problem: managing files over time

- Have you ever had files like this?
 - `thesis_v1.docx`
 - `thesis_v2_final.docx`
 - `thesis_v2_final_REAL.docx`
 - `thesis_v2_final_REAL_submitted.docx`
- What changed between versions?
- Which version has the correct analysis?

The problem: managing files over time

- Have you ever had files like this?
 - `thesis_v1.docx`
 - `thesis_v2_final.docx`
 - `thesis_v2_final_REAL.docx`
 - `thesis_v2_final_REAL_submitted.docx`
- What changed between versions?
- Which version has the correct analysis?
- How do you collaborate without overwriting each other's work?

Version control: a better way

Version control is a system that records changes to files over time

- One file, complete history

Version control: a better way

Version control is a system that records changes to files over time

- One file, complete history
- Every change is recorded with a description

Version control: a better way

Version control is a system that records changes to files over time

- One file, complete history
- Every change is recorded with a description
- Can go back to any previous state

Version control: a better way

Version control is a system that records changes to files over time

- One file, complete history
- Every change is recorded with a description
- Can go back to any previous state
- Multiple people can work simultaneously

Why version control for research?

- **Reproducibility:** Track exactly what you did and when

Why version control for research?

- **Reproducibility:** Track exactly what you did and when
- **Backup:** Your work is safely stored, even if your laptop dies

Why version control for research?

- **Reproducibility:** Track exactly what you did and when
- **Backup:** Your work is safely stored, even if your laptop dies
- **Collaboration:** Work with others without email chains of files

Why version control for research?

- **Reproducibility:** Track exactly what you did and when
- **Backup:** Your work is safely stored, even if your laptop dies
- **Collaboration:** Work with others without email chains of files
- **Transparency:** Share your code with the research community

Why version control for research?

- **Reproducibility:** Track exactly what you did and when
- **Backup:** Your work is safely stored, even if your laptop dies
- **Collaboration:** Work with others without email chains of files
- **Transparency:** Share your code with the research community

Why version control for research?

- **Reproducibility:** Track exactly what you did and when
- **Backup:** Your work is safely stored, even if your laptop dies
- **Collaboration:** Work with others without email chains of files
- **Transparency:** Share your code with the research community
- Many journals now require or encourage sharing code via GitHub

Git and GitHub

Git

- A version control system
- Runs locally on your computer
- Tracks changes to files

GitHub

- A web platform that hosts Git repositories
- Stores your code online
- Enables sharing and collaboration

The basic Git workflow

1. **Make changes** to your files (write code, edit text)

The basic Git workflow

1. **Make changes** to your files (write code, edit text)
2. **Stage** the changes you want to save

The basic Git workflow

1. **Make changes** to your files (write code, edit text)
2. **Stage** the changes you want to save
 - “These are the files I want to include in my next snapshot”

The basic Git workflow

1. **Make changes** to your files (write code, edit text)
2. **Stage** the changes you want to save
 - “These are the files I want to include in my next snapshot”
3. **Commit** the staged changes with a message

The basic Git workflow

1. **Make changes** to your files (write code, edit text)
2. **Stage** the changes you want to save
 - “These are the files I want to include in my next snapshot”
3. **Commit** the staged changes with a message
 - A snapshot of your project at this moment

The basic Git workflow

1. **Make changes** to your files (write code, edit text)
2. **Stage** the changes you want to save
 - “These are the files I want to include in my next snapshot”
3. **Commit** the staged changes with a message
 - A snapshot of your project at this moment
4. **Push** your commits to GitHub

The basic Git workflow

1. **Make changes** to your files (write code, edit text)
2. **Stage** the changes you want to save
 - “These are the files I want to include in my next snapshot”
3. **Commit** the staged changes with a message
 - A snapshot of your project at this moment
4. **Push** your commits to GitHub
 - Upload your local changes to the cloud

Ways to use Git

- **GitHub web interface:** Create repos, upload files, edit directly

Ways to use Git

- **GitHub web interface:** Create repos, upload files, edit directly
 - Simple but limited

Ways to use Git

- **GitHub web interface:** Create repos, upload files, edit directly
 - Simple but limited
- **Command line:** Most powerful and flexible

Ways to use Git

- **GitHub web interface:** Create repos, upload files, edit directly
 - Simple but limited
- **Command line:** Most powerful and flexible
 - `git add`, `git commit`, `git push`

Ways to use Git

- **GitHub web interface:** Create repos, upload files, edit directly
 - Simple but limited
- **Command line:** Most powerful and flexible
 - `git add`, `git commit`, `git push`
- **RStudio:** Built-in Git integration

Ways to use Git

- **GitHub web interface:** Create repos, upload files, edit directly
 - Simple but limited
- **Command line:** Most powerful and flexible
 - `git add`, `git commit`, `git push`
- **RStudio:** Built-in Git integration
 - Point-and-click interface

Ways to use Git

- **GitHub web interface:** Create repos, upload files, edit directly
 - Simple but limited
- **Command line:** Most powerful and flexible
 - `git add`, `git commit`, `git push`
- **RStudio:** Built-in Git integration
 - Point-and-click interface

Ways to use Git

- **GitHub web interface:** Create repos, upload files, edit directly
 - Simple but limited
- **Command line:** Most powerful and flexible
 - `git add`, `git commit`, `git push`
- **RStudio:** Built-in Git integration
 - Point-and-click interface
- All do the same thing—choose what works for you

Assignment 1

- Create a GitHub account (if you don't have one)

Assignment 1

- Create a GitHub account (if you don't have one)
- Create a **public** repository for this course

Assignment 1

- Create a GitHub account (if you don't have one)
- Create a **public** repository for this course
- Set up your README and folder structure

Assignment 1

- Create a GitHub account (if you don't have one)
- Create a **public** repository for this course
- Set up your README and folder structure
- Create a simple .R file

Assignment 1

- Create a GitHub account (if you don't have one)
- Create a **public** repository for this course
- Set up your README and folder structure
- Create a simple .R file

Assignment 1

- Create a GitHub account (if you don't have one)
 - Create a **public** repository for this course
 - Set up your README and folder structure
 - Create a simple .R file
-
- This repository is where you'll submit all your assignments

Assignment 1

- Create a GitHub account (if you don't have one)
 - Create a **public** repository for this course
 - Set up your README and folder structure
 - Create a simple .R file
-
- This repository is where you'll submit all your assignments
 - Detailed instructions in the assignment document

What makes a good analysis?

- Clear research question

What makes a good analysis?

- Clear research question
- Appropriate data for the question

What makes a good analysis?

- Clear research question
- Appropriate data for the question
- Right statistical model for the data

What makes a good analysis?

- Clear research question
- Appropriate data for the question
- Right statistical model for the data
- Correct interpretation of results

What makes a good analysis?

- Clear research question
- Appropriate data for the question
- Right statistical model for the data
- Correct interpretation of results
- Honest about limitations and uncertainty

Looking ahead

- Next session: Applied regression

Looking ahead

- Next session: Applied regression
- Regression as conditional expectations

Looking ahead

- Next session: Applied regression
- Regression as conditional expectations
- Multiple regression and control variables

Looking ahead

- Next session: Applied regression
- Regression as conditional expectations
- Multiple regression and control variables
- Interaction effects and presenting results

For next week

- Check readings if needed
- Review your notes on OLS from AQMSS-I
- **Finish Assignment 1**

Questions?