

Review: Research Design & Advanced Topics

Francisco Villamil

Applied Quantitative Methods II
MA in Social Sciences, Spring 2026

Today's second half

- The course has been about **estimation**: how to fit models

Today's second half

- The course has been about **estimation**: how to fit models
- But fitting a model is rarely the hard part of empirical work

Today's second half

- The course has been about **estimation**: how to fit models
- But fitting a model is rarely the hard part of empirical work
- Today: what you actually need to answer a research question

Today's second half

- The course has been about **estimation**: how to fit models
- But fitting a model is rarely the hard part of empirical work
- Today: what you actually need to answer a research question
 - Research design and identification

Today's second half

- The course has been about **estimation**: how to fit models
- But fitting a model is rarely the hard part of empirical work
- Today: what you actually need to answer a research question
 - Research design and identification
 - Causal inference methods beyond FE/DiD

Today's second half

- The course has been about **estimation**: how to fit models
- But fitting a model is rarely the hard part of empirical work
- Today: what you actually need to answer a research question
 - Research design and identification
 - Causal inference methods beyond FE/DiD
 - Defending your results: robustness and validity

Today's second half

- The course has been about **estimation**: how to fit models
- But fitting a model is rarely the hard part of empirical work
- Today: what you actually need to answer a research question
 - Research design and identification
 - Causal inference methods beyond FE/DiD
 - Defending your results: robustness and validity
 - Common pitfalls when applying all of this

Today's second half

- The course has been about **estimation**: how to fit models
- But fitting a model is rarely the hard part of empirical work
- Today: what you actually need to answer a research question
 - Research design and identification
 - Causal inference methods beyond FE/DiD
 - Defending your results: robustness and validity
 - Common pitfalls when applying all of this
- Plus a pointer to two advanced topics (not covered today)

Roadmap

From estimation to identification

Causal inference toolkit

Defending a causal claim

Choosing the right method & pitfalls

What you have learned

- Estimation tools for many kinds of outcomes

What you have learned

- Estimation tools for many kinds of outcomes
 - OLS, logit, count / ordinal / duration models

What you have learned

- Estimation tools for many kinds of outcomes
 - OLS, logit, count / ordinal / duration models
- Designs that exploit data structure

What you have learned

- Estimation tools for many kinds of outcomes
 - OLS, logit, count / ordinal / duration models
- Designs that exploit data structure
 - Fixed effects, two-way FE, DiD (including staggered)

What you have learned

- Estimation tools for many kinds of outcomes
 - OLS, logit, count / ordinal / duration models
- Designs that exploit data structure
 - Fixed effects, two-way FE, DiD (including staggered)
 - Spatial models

What you have learned

- Estimation tools for many kinds of outcomes
 - OLS, logit, count / ordinal / duration models
- Designs that exploit data structure
 - Fixed effects, two-way FE, DiD (including staggered)
 - Spatial models
- How to interpret and present results

What you have learned

- Estimation tools for many kinds of outcomes
 - OLS, logit, count / ordinal / duration models
- Designs that exploit data structure
 - Fixed effects, two-way FE, DiD (including staggered)
 - Spatial models
- How to interpret and present results
 - Predicted values, marginal effects, coefficient plots

What you have learned

- Estimation tools for many kinds of outcomes
 - OLS, logit, count / ordinal / duration models
- Designs that exploit data structure
 - Fixed effects, two-way FE, DiD (including staggered)
 - Spatial models
- How to interpret and present results
 - Predicted values, marginal effects, coefficient plots
- Plus computing practices to make all of this reproducible

Two different questions

Estimation: given a model, what are the parameters?

Identification: do those parameters mean what I think they mean?

Two different questions

Estimation: given a model, what are the parameters?

Identification: do those parameters mean what I think they mean?

Your R output answers the first.

Only your research design can answer the second.

The fundamental problem of causal inference

- A **causal effect** is the difference between what *did* happen and what *would have* happened otherwise

The fundamental problem of causal inference

- A **causal effect** is the difference between what *did* happen and what *would have* happened otherwise
- For unit i with treatment D :

$$\tau_i = Y_i(1) - Y_i(0)$$

The fundamental problem of causal inference

- A **causal effect** is the difference between what *did* happen and what *would have* happened otherwise
- For unit i with treatment D :

$$\tau_i = Y_i(1) - Y_i(0)$$

- We only ever observe **one** of these two worlds

The fundamental problem of causal inference

- A **causal effect** is the difference between what *did* happen and what *would have* happened otherwise
- For unit i with treatment D :

$$\tau_i = Y_i(1) - Y_i(0)$$

- We only ever observe **one** of these two worlds
 - If i was treated, $Y_i(1)$ is observed and $Y_i(0)$ is missing

The fundamental problem of causal inference

- A **causal effect** is the difference between what *did* happen and what *would have* happened otherwise
- For unit i with treatment D :

$$\tau_i = Y_i(1) - Y_i(0)$$

- We only ever observe **one** of these two worlds
 - If i was treated, $Y_i(1)$ is observed and $Y_i(0)$ is missing
 - If i was not treated, the opposite

The fundamental problem of causal inference

- A **causal effect** is the difference between what *did* happen and what *would have* happened otherwise
- For unit i with treatment D :

$$\tau_i = Y_i(1) - Y_i(0)$$

- We only ever observe **one** of these two worlds
 - If i was treated, $Y_i(1)$ is observed and $Y_i(0)$ is missing
 - If i was not treated, the opposite
- Causal inference = principled way to estimate the missing counterfactual

What we can estimate on average

- We can't recover τ_i for individuals
- But we can target averages:
 - **ATE** — average over everyone: $E[Y(1) - Y(0)]$
 - **ATT** — average over the treated: $E[Y(1) - Y(0) | D = 1]$
 - **LATE** — average over those who *respond* to a treatment (IV, RDD)
- Different methods estimate different quantities
 - When you compare two papers: are they even targeting the same thing?

Causal models and DAGs

- To approximate the counterfactual we need a model of *why* units are treated in the first place

Causal models and DAGs

- To approximate the counterfactual we need a model of *why* units are treated in the first place
- Directed Acyclic Graphs (DAGs): nodes are variables, arrows are direct causal effects

Causal models and DAGs

- To approximate the counterfactual we need a model of *why* units are treated in the first place
- Directed Acyclic Graphs (DAGs): nodes are variables, arrows are direct causal effects
- Three structures you must recognize:

Causal models and DAGs

- To approximate the counterfactual we need a model of *why* units are treated in the first place
- Directed Acyclic Graphs (DAGs): nodes are variables, arrows are direct causal effects
- Three structures you must recognize:
 - **Confounder** $D \leftarrow Z \rightarrow Y$

Causal models and DAGs

- To approximate the counterfactual we need a model of *why* units are treated in the first place
- Directed Acyclic Graphs (DAGs): nodes are variables, arrows are direct causal effects
- Three structures you must recognize:
 - **Confounder** $D \leftarrow Z \rightarrow Y$
 - **Mediator** $D \rightarrow M \rightarrow Y$


Causal models and DAGs

- To approximate the counterfactual we need a model of *why* units are treated in the first place
- Directed Acyclic Graphs (DAGs): nodes are variables, arrows are direct causal effects
- Three structures you must recognize:
 - **Confounder** $D \leftarrow Z \rightarrow Y$
 - **Mediator** $D \rightarrow M \rightarrow Y$
 - **Collider** $D \rightarrow C \leftarrow Y$

Causal models and DAGs


- To approximate the counterfactual we need a model of *why* units are treated in the first place
- Directed Acyclic Graphs (DAGs): nodes are variables, arrows are direct causal effects
- Three structures you must recognize:
 - **Confounder** $D \leftarrow Z \rightarrow Y$
 - **Mediator** $D \rightarrow M \rightarrow Y$
 - **Collider** $D \rightarrow C \leftarrow Y$
- What you control for — and what you do **not** — depends on which structure you are looking at

Confounding: the standard case



- Z affects both D and Y : a back-door path
- Un-adjusted comparison of treated vs. untreated mixes two effects
- **Solution:** condition on Z (or use a design that blocks the back door)

Colliders: do not control!



- C is caused *by* both D and Y
- Conditioning on C creates a spurious association between D and Y
- Classic: “smart people lack social skills” (if you condition on being hired)

Post-treatment bias

- Special case of collider bias that is very easy to miss

Post-treatment bias

- Special case of collider bias that is very easy to miss
- A **post-treatment** variable is one that is itself affected by the treatment

Post-treatment bias

- Special case of collider bias that is very easy to miss
- A **post-treatment** variable is one that is itself affected by the treatment
- Controlling for it:

Post-treatment bias

- Special case of collider bias that is very easy to miss
- A **post-treatment** variable is one that is itself affected by the treatment
- Controlling for it:
 - blocks part of the mechanism (if it is a mediator), **or**

Post-treatment bias

- Special case of collider bias that is very easy to miss
- A **post-treatment** variable is one that is itself affected by the treatment
- Controlling for it:
 - blocks part of the mechanism (if it is a mediator), **or**
 - opens a collider path (if it is also affected by other unobservables)

Post-treatment bias

- Special case of collider bias that is very easy to miss
- A **post-treatment** variable is one that is itself affected by the treatment
- Controlling for it:
 - blocks part of the mechanism (if it is a mediator), **or**
 - opens a collider path (if it is also affected by other unobservables)
- Rule of thumb: if in doubt about whether a variable is pre- or post-treatment, leave it out

So what does your regression assume?

Every regression coefficient you report implicitly assumes:

1. You have **closed all back doors** (no omitted confounders)
2. You have **not opened new ones** (no controlling for colliders or mediators)
3. The functional form is roughly right (linearity, additivity)

So what does your regression assume?

Every regression coefficient you report implicitly assumes:

1. You have **closed all back doors** (no omitted confounders)
2. You have **not opened new ones** (no controlling for colliders or mediators)
3. The functional form is roughly right (linearity, additivity)

(1) and (2) are **design** assumptions, not something you can test.

Roadmap

From estimation to identification

Causal inference toolkit

Defending a causal claim

Choosing the right method & pitfalls

What you already have

Two of the four big causal-inference designs are already in your toolkit:

- **Fixed effects** — control for all time-invariant unobservables at some level (unit, region, year. . .)
- **Difference-in-differences** — use a never-treated comparison group plus pre/post variation; identifies the ATT under parallel trends

Two more that you should at least recognize:

- **Regression discontinuity design (RDD)**
- **Instrumental variables (IV)**

Estimating DiD

- DiD is not a special estimator — it's just OLS on a 2×2 interaction

Estimating DiD

- DiD is not a special estimator — it's just OLS on a 2×2 interaction

$$Y_{it} = \alpha + \beta_1 \text{Treat}_i + \beta_2 \text{Post}_t + \beta_3 (\text{Treat}_i \times \text{Post}_t) + \varepsilon_{it}$$

Estimating DiD

- DiD is not a special estimator — it's just OLS on a 2×2 interaction

$$Y_{it} = \alpha + \beta_1 \text{Treat}_i + \beta_2 \text{Post}_t + \beta_3 (\text{Treat}_i \times \text{Post}_t) + \varepsilon_{it}$$

- β_3 is the DiD estimate

Estimating DiD

- DiD is not a special estimator — it's just OLS on a 2×2 interaction

$$Y_{it} = \alpha + \beta_1 \text{Treat}_i + \beta_2 \text{Post}_t + \beta_3 (\text{Treat}_i \times \text{Post}_t) + \varepsilon_{it}$$

- β_3 is the DiD estimate
- In practice: two-way fixed effects absorb Treat_i and Post_t

Estimating DiD

- DiD is not a special estimator — it's just OLS on a 2×2 interaction

$$Y_{it} = \alpha + \beta_1 \text{Treat}_i + \beta_2 \text{Post}_t + \beta_3 (\text{Treat}_i \times \text{Post}_t) + \varepsilon_{it}$$

- β_3 is the DiD estimate
- In practice: two-way fixed effects absorb Treat_i and Post_t
`feols(y ~ treat_post | unit + time, data = d)`

Estimating DiD

- DiD is not a special estimator — it's just OLS on a 2×2 interaction

$$Y_{it} = \alpha + \beta_1 \text{Treat}_i + \beta_2 \text{Post}_t + \beta_3 (\text{Treat}_i \times \text{Post}_t) + \varepsilon_{it}$$

- β_3 is the DiD estimate
- In practice: two-way fixed effects absorb Treat_i and Post_t
`feols(y ~ treat_post | unit + time, data = d)`
- For staggered adoption: use modern estimators
(Callaway–Sant’Anna, Sun–Abraham) — see Session 6

Regression discontinuity design

- **Setup:** treatment assignment depends on whether a *running variable* crosses a cutoff

Regression discontinuity design

- **Setup:** treatment assignment depends on whether a *running variable* crosses a cutoff
 - Incumbency advantage: winners of close elections

Regression discontinuity design

- **Setup:** treatment assignment depends on whether a *running variable* crosses a cutoff
 - Incumbency advantage: winners of close elections
 - Scholarships: eligibility based on test score cutoff

Regression discontinuity design

- **Setup:** treatment assignment depends on whether a *running variable* crosses a cutoff
 - Incumbency advantage: winners of close elections
 - Scholarships: eligibility based on test score cutoff
 - Conscription: draft lottery numbers below a threshold


Regression discontinuity design

- **Setup:** treatment assignment depends on whether a *running variable* crosses a cutoff
 - Incumbency advantage: winners of close elections
 - Scholarships: eligibility based on test score cutoff
 - Conscription: draft lottery numbers below a threshold
- **Idea:** units just above and just below the cutoff should be similar in everything *except* the treatment

Regression discontinuity design

- **Setup:** treatment assignment depends on whether a *running variable* crosses a cutoff
 - Incumbency advantage: winners of close elections
 - Scholarships: eligibility based on test score cutoff
 - Conscription: draft lottery numbers below a threshold
- **Idea:** units just above and just below the cutoff should be similar in everything *except* the treatment
- Estimates a **LATE**: effect for units near the threshold

RDD visually



- The jump at the cutoff is the treatment effect
- The trend on each side controls for whatever else varies smoothly

Estimating RDD

- Center the running variable at the cutoff: $\tilde{X}_i = X_i - c$

Estimating RDD

- Center the running variable at the cutoff: $\tilde{X}_i = X_i - c$
- Define treatment: $D_i = \mathbf{1}(\tilde{X}_i \geq 0)$

Estimating RDD

- Center the running variable at the cutoff: $\tilde{X}_i = X_i - c$
- Define treatment: $D_i = \mathbf{1}(\tilde{X}_i \geq 0)$
- Run a local linear regression with slopes allowed to differ on each side:

Estimating RDD

- Center the running variable at the cutoff: $\tilde{X}_i = X_i - c$
- Define treatment: $D_i = \mathbf{1}(\tilde{X}_i \geq 0)$
- Run a local linear regression with slopes allowed to differ on each side:

$$Y_i = \alpha + \tau D_i + \beta_1 \tilde{X}_i + \beta_2 (D_i \times \tilde{X}_i) + \varepsilon_i$$

Estimating RDD

- Center the running variable at the cutoff: $\tilde{X}_i = X_i - c$
- Define treatment: $D_i = \mathbf{1}(\tilde{X}_i \geq 0)$
- Run a local linear regression with slopes allowed to differ on each side:

$$Y_i = \alpha + \tau D_i + \beta_1 \tilde{X}_i + \beta_2 (D_i \times \tilde{X}_i) + \varepsilon_i$$

- τ is the RDD estimate (the jump at the cutoff)

Estimating RDD

- Center the running variable at the cutoff: $\tilde{X}_i = X_i - c$
- Define treatment: $D_i = \mathbf{1}(\tilde{X}_i \geq 0)$
- Run a local linear regression with slopes allowed to differ on each side:

$$Y_i = \alpha + \tau D_i + \beta_1 \tilde{X}_i + \beta_2 (D_i \times \tilde{X}_i) + \varepsilon_i$$

- τ is the RDD estimate (the jump at the cutoff)
- Key choice: **bandwidth** — how wide a window around the cutoff

Estimating RDD

- Center the running variable at the cutoff: $\tilde{X}_i = X_i - c$
- Define treatment: $D_i = \mathbf{1}(\tilde{X}_i \geq 0)$
- Run a local linear regression with slopes allowed to differ on each side:

$$Y_i = \alpha + \tau D_i + \beta_1 \tilde{X}_i + \beta_2 (D_i \times \tilde{X}_i) + \varepsilon_i$$

- τ is the RDD estimate (the jump at the cutoff)
- Key choice: **bandwidth** — how wide a window around the cutoff
 - Data-driven: Calonico, Cattaneo & Titiunik

Estimating RDD

- Center the running variable at the cutoff: $\tilde{X}_i = X_i - c$
- Define treatment: $D_i = \mathbf{1}(\tilde{X}_i \geq 0)$
- Run a local linear regression with slopes allowed to differ on each side:

$$Y_i = \alpha + \tau D_i + \beta_1 \tilde{X}_i + \beta_2 (D_i \times \tilde{X}_i) + \varepsilon_i$$

- τ is the RDD estimate (the jump at the cutoff)
- Key choice: **bandwidth** — how wide a window around the cutoff
 - Data-driven: Calonico, Cattaneo & Titiunik
- R: `rdrobust::rdrobust(y, x, c = 0)`

Instrumental variables



- Find a variable Z that moves D but has no other path to Y

Instrumental variables



- Find a variable Z that moves D but has no other path to Y
- Use only the variation in D explained by Z

Instrumental variables



- Find a variable Z that moves D but has no other path to Y
- Use only the variation in D explained by Z
- Two assumptions:

Instrumental variables



- Find a variable Z that moves D but has no other path to Y
- Use only the variation in D explained by Z
- Two assumptions:
 - **Relevance:** Z actually affects D (testable)

Instrumental variables



- Find a variable Z that moves D but has no other path to Y
- Use only the variation in D explained by Z
- Two assumptions:
 - **Relevance**: Z actually affects D (testable)
 - **Exclusion restriction**: Z affects Y only through D (not testable)

Estimating IV: two-stage least squares

- **First stage:** regress the treatment on the instrument

Estimating IV: two-stage least squares

- **First stage:** regress the treatment on the instrument

$$D_i = \pi_0 + \pi_1 Z_i + u_i \Rightarrow \hat{D}_i$$

Estimating IV: two-stage least squares

- **First stage:** regress the treatment on the instrument

$$D_i = \pi_0 + \pi_1 Z_i + u_i \Rightarrow \hat{D}_i$$

- **Second stage:** regress the outcome on the *predicted* treatment

Estimating IV: two-stage least squares

- **First stage:** regress the treatment on the instrument

$$D_i = \pi_0 + \pi_1 Z_i + u_i \Rightarrow \hat{D}_i$$

- **Second stage:** regress the outcome on the *predicted* treatment

$$Y_i = \alpha + \beta \hat{D}_i + \varepsilon_i$$

Estimating IV: two-stage least squares

- **First stage:** regress the treatment on the instrument

$$D_i = \pi_0 + \pi_1 Z_i + u_i \Rightarrow \hat{D}_i$$

- **Second stage:** regress the outcome on the *predicted* treatment

$$Y_i = \alpha + \beta \hat{D}_i + \varepsilon_i$$

- β is the IV estimate: a LATE for units whose D responds to Z

Estimating IV: two-stage least squares

- **First stage:** regress the treatment on the instrument

$$D_i = \pi_0 + \pi_1 Z_i + u_i \Rightarrow \hat{D}_i$$

- **Second stage:** regress the outcome on the *predicted* treatment

$$Y_i = \alpha + \beta \hat{D}_i + \varepsilon_i$$

- β is the IV estimate: a LATE for units whose D responds to Z
- Always report the first stage, and an F-statistic on π_1

Estimating IV: two-stage least squares

- **First stage:** regress the treatment on the instrument

$$D_i = \pi_0 + \pi_1 Z_i + u_i \Rightarrow \hat{D}_i$$

- **Second stage:** regress the outcome on the *predicted* treatment

$$Y_i = \alpha + \beta \hat{D}_i + \varepsilon_i$$

- β is the IV estimate: a LATE for units whose D responds to Z
- Always report the first stage, and an **F-statistic** on π_1
 - Rule of thumb: $F < 10$ signals a weak instrument

Estimating IV: two-stage least squares

- **First stage:** regress the treatment on the instrument

$$D_i = \pi_0 + \pi_1 Z_i + u_i \Rightarrow \hat{D}_i$$

- **Second stage:** regress the outcome on the *predicted* treatment

$$Y_i = \alpha + \beta \hat{D}_i + \varepsilon_i$$

- β is the IV estimate: a LATE for units whose D responds to Z
- Always report the first stage, and an **F-statistic** on π_1
 - Rule of thumb: $F < 10$ signals a weak instrument
- R: `feols(y ~ 1 | d ~ z)` or `AER::ivreg()`

The four design templates

Design	Exogenous variation	Key assumption
FE	Within-unit differences	No time-varying confounders
DiD	Treated vs. comparison, pre/post	Parallel trends
RDD	Jump at cutoff	No manipulation at cutoff
IV	Instrument $\rightarrow$ treatment	Exclusion restriction

- You do not pick a design, *the data* pick one for you
- What source of exogenous variation does the world hand you?

Matching: complement, not substitute

- Not a design in the same sense — but a useful preprocessing step

Matching: complement, not substitute

- Not a design in the same sense — but a useful preprocessing step
- **Idea:** pair each treated unit with one or more untreated units that look similar on observed covariates

Matching: complement, not substitute

- Not a design in the same sense — but a useful preprocessing step
- **Idea:** pair each treated unit with one or more untreated units that look similar on observed covariates
- Compare only within matched pairs

Matching: complement, not substitute

- Not a design in the same sense — but a useful preprocessing step
- **Idea:** pair each treated unit with one or more untreated units that look similar on observed covariates
- Compare only within matched pairs
- What it fixes: imbalance in observed covariates and reliance on linear-in-parameters extrapolation

Matching: complement, not substitute

- Not a design in the same sense — but a useful preprocessing step
- **Idea:** pair each treated unit with one or more untreated units that look similar on observed covariates
- Compare only within matched pairs
- What it fixes: imbalance in observed covariates and reliance on linear-in-parameters extrapolation
- What it does **not** fix: omitted confounders, selection on unobservables

Matching: complement, not substitute

- Not a design in the same sense — but a useful preprocessing step
- **Idea:** pair each treated unit with one or more untreated units that look similar on observed covariates
- Compare only within matched pairs
- What it fixes: imbalance in observed covariates and reliance on linear-in-parameters extrapolation
- What it does **not** fix: omitted confounders, selection on unobservables
- Often used *with* regression or DiD, not instead of

Roadmap

From estimation to identification

Causal inference toolkit

Defending a causal claim

Choosing the right method & pitfalls

Your main result is the beginning

- You have a point estimate, a standard error, a figure
- A skeptical reader should be asking:
 - What if you had specified it differently?
 - What if your identifying assumption is wrong?
 - What if you are measuring the wrong thing?
 - Would the result generalize elsewhere?
- Your job is to answer those questions **before they do**

Robustness: alternative specifications

- Change one thing at a time, show it does not matter much

Robustness: alternative specifications

- Change one thing at a time, show it does not matter much
- Measurement of Y : alternative indicators, different scales

Robustness: alternative specifications

- Change one thing at a time, show it does not matter much
- Measurement of Y : alternative indicators, different scales
- Measurement of D : dichotomize vs. continuous, source

Robustness: alternative specifications

- Change one thing at a time, show it does not matter much
- Measurement of Y : alternative indicators, different scales
- Measurement of D : dichotomize vs. continuous, source
- Sample: drop influential units, drop specific sub-groups

Robustness: alternative specifications

- Change one thing at a time, show it does not matter much
- Measurement of Y : alternative indicators, different scales
- Measurement of D : dichotomize vs. continuous, source
- Sample: drop influential units, drop specific sub-groups
- Model: add/remove controls, change FE structure, change SEs

Robustness: alternative specifications

- Change one thing at a time, show it does not matter much
- Measurement of Y : alternative indicators, different scales
- Measurement of D : dichotomize vs. continuous, source
- Sample: drop influential units, drop specific sub-groups
- Model: add/remove controls, change FE structure, change SEs
- Functional form: log, quadratic, non-parametric

Robustness: alternative specifications

- Change one thing at a time, show it does not matter much
- Measurement of Y : alternative indicators, different scales
- Measurement of D : dichotomize vs. continuous, source
- Sample: drop influential units, drop specific sub-groups
- Model: add/remove controls, change FE structure, change SEs
- Functional form: log, quadratic, non-parametric
- Display as a coefficient plot across specifications, not a 10-column table no one reads

Placebo tests

- Look for an effect *where there should be none*

Placebo tests

- Look for an effect *where there should be none*
- If you find it, something is wrong with your design

Placebo tests

- Look for an effect *where there should be none*
- If you find it, something is wrong with your design
- Three flavors (Eggers, Tuñón & Dafoe 2023):

Placebo tests

- Look for an effect *where there should be none*
- If you find it, something is wrong with your design
- Three flavors (Eggers, Tuñón & Dafoe 2023):
 - **Placebo outcome**: different Y that the treatment should not affect

Placebo tests

- Look for an effect *where there should be none*
- If you find it, something is wrong with your design
- Three flavors (Eggers, Tuñón & Dafoe 2023):
 - **Placebo outcome**: different Y that the treatment should not affect
 - **Placebo treatment**: a fake D that should have no effect

Placebo tests

- Look for an effect *where there should be none*
- If you find it, something is wrong with your design
- Three flavors (Eggers, Tuñón & Dafoe 2023):
 - **Placebo outcome**: different Y that the treatment should not affect
 - **Placebo treatment**: a fake D that should have no effect
 - **Placebo population**: a subgroup not exposed to D

Placebo tests in DiD

- Most important placebo in a DiD setup: **pre-trends**

Placebo tests in DiD

- Most important placebo in a DiD setup: **pre-trends**
- Before the treatment, the two groups should move in parallel

Placebo tests in DiD

- Most important placebo in a DiD setup: **pre-trends**
- Before the treatment, the two groups should move in parallel
- If pre-treatment "leads" are significant, parallel trends is suspect

Placebo tests in DiD

- Most important placebo in a DiD setup: **pre-trends**
- Before the treatment, the two groups should move in parallel
- If pre-treatment "leads" are significant, parallel trends is suspect
- Event-study plots show this visually:

Placebo tests in DiD

- Most important placebo in a DiD setup: **pre-trends**
- Before the treatment, the two groups should move in parallel
- If pre-treatment "leads" are significant, parallel trends is suspect
- Event-study plots show this visually:
 - Coefficients on event-time dummies

Placebo tests in DiD

- Most important placebo in a DiD setup: **pre-trends**
- Before the treatment, the two groups should move in parallel
- If pre-treatment "leads" are significant, parallel trends is suspect
- Event-study plots show this visually:
 - Coefficients on event-time dummies
 - Pre-period coefficients should be flat and near zero

Placebo tests in DiD

- Most important placebo in a DiD setup: **pre-trends**
- Before the treatment, the two groups should move in parallel
- If pre-treatment "leads" are significant, parallel trends is suspect
- Event-study plots show this visually:
 - Coefficients on event-time dummies
 - Pre-period coefficients should be flat and near zero
 - Post-period coefficients show the dynamic effect

Additional implications

- If your theory is correct, what *else* should be true in the data?

Additional implications

- If your theory is correct, what *else* should be true in the data?
- **Mechanism checks:** intermediate variables your theory predicts

Additional implications

- If your theory is correct, what *else* should be true in the data?
- **Mechanism checks:** intermediate variables your theory predicts
- **Heterogeneity:** groups where the effect should be larger or smaller

Additional implications

- If your theory is correct, what *else* should be true in the data?
- **Mechanism checks:** intermediate variables your theory predicts
- **Heterogeneity:** groups where the effect should be larger or smaller
- **Rule-out tests:** observable patterns that competing explanations would imply but yours would not

Additional implications

- If your theory is correct, what *else* should be true in the data?
- **Mechanism checks**: intermediate variables your theory predicts
- **Heterogeneity**: groups where the effect should be larger or smaller
- **Rule-out tests**: observable patterns that competing explanations would imply but yours would not
- A paper that only shows the main effect is weaker than one that shows the main effect **plus** two or three such auxiliary findings

Measurement and construct validity

- Are you measuring what you think you are measuring?

Measurement and construct validity

- Are you measuring what you think you are measuring?
- **Treatment:** when you say $D = 1$, what is actually being delivered?

Measurement and construct validity

- Are you measuring what you think you are measuring?
- **Treatment:** when you say $D = 1$, what is actually being delivered?
 - Bundled treatments (media exposure = content + format + novelty...)

Measurement and construct validity

- Are you measuring what you think you are measuring?
- **Treatment:** when you say $D = 1$, what is actually being delivered?
 - Bundled treatments (media exposure = content + format + novelty...)
- **Outcome:** does the indicator really capture the concept?

Measurement and construct validity

- Are you measuring what you think you are measuring?
- **Treatment:** when you say $D = 1$, what is actually being delivered?
 - Bundled treatments (media exposure = content + format + novelty...)
- **Outcome:** does the indicator really capture the concept?
 - Polarization = partisan thermometer gap? policy-view distance? affective feelings?

Measurement and construct validity

- Are you measuring what you think you are measuring?
- **Treatment:** when you say $D = 1$, what is actually being delivered?
 - Bundled treatments (media exposure = content + format + novelty...)
- **Outcome:** does the indicator really capture the concept?
 - Polarization = partisan thermometer gap? policy-view distance? affective feelings?
- Identification can be perfect and the claim still be wrong if the measurement doesn't line up with the theory

External validity

- Under what *conditions* does the effect hold?

External validity

- Under what *conditions* does the effect hold?
- Four dimensions (Egami & Hartman 2023):

External validity

- Under what *conditions* does the effect hold?
- Four dimensions (Egami & Hartman 2023):
 - **Populations**: units beyond your sample

External validity

- Under what *conditions* does the effect hold?
- Four dimensions (Egami & Hartman 2023):
 - **Populations**: units beyond your sample
 - **Treatments**: variants of the intervention

External validity

- Under what *conditions* does the effect hold?
- Four dimensions (Egami & Hartman 2023):
 - **Populations**: units beyond your sample
 - **Treatments**: variants of the intervention
 - **Outcomes**: related but different Y s

External validity

- Under what *conditions* does the effect hold?
- Four dimensions (Egami & Hartman 2023):
 - **Populations**: units beyond your sample
 - **Treatments**: variants of the intervention
 - **Outcomes**: related but different Ys
 - **Contexts**: other settings, places, times

External validity

- Under what *conditions* does the effect hold?
- Four dimensions (Egami & Hartman 2023):
 - **Populations**: units beyond your sample
 - **Treatments**: variants of the intervention
 - **Outcomes**: related but different Ys
 - **Contexts**: other settings, places, times
- A clean estimate in one sample does not automatically generalize

External validity

- Under what *conditions* does the effect hold?
- Four dimensions (Egami & Hartman 2023):
 - **Populations**: units beyond your sample
 - **Treatments**: variants of the intervention
 - **Outcomes**: related but different Ys
 - **Contexts**: other settings, places, times
- A clean estimate in one sample does not automatically generalize
- **Temporal validity** (Munger 2023): does it still hold 10 years from now?

Roadmap

From estimation to identification

Causal inference toolkit

Defending a causal claim

Choosing the right method & pitfalls

Starting point: your question type

Goal	Tools
Describe a pattern	Summary stats, visualization, simple regression
Predict Y from X	Regression, ML (cross-validated)
Explain Y causally	Design-based methods: FE, DiD, RDD, IV, experiments
Generalize a mechanism	Theory + replication across contexts

- Most confusion in applied work comes from mixing these up
 - e.g. treating predictive coefficients as causal

Starting point: your outcome

Outcome type	Method
Continuous	OLS (always a good first pass)
Binary	Logit, probit (or LPM for interpretability)
Count (non-negative integer)	Poisson, negative binomial
Ordered categories	Ordered logit / probit
Time-to-event	Cox, Kaplan–Meier, parametric survival
Panel data	+ fixed effects, + clustered SEs
Spatial data	+ spatial weights, + Moran's I

- This is a cheat-sheet for your final essay — keep it

Pitfall 1: p-hacking and forking paths

- Garden of forking paths (Gelman & Loken 2013)

Pitfall 1: p-hacking and forking paths

- Garden of forking paths (Gelman & Loken 2013)
- Every analysis involves many small choices:

Pitfall 1: p-hacking and forking paths

- Garden of forking paths (Gelman & Loken 2013)
- Every analysis involves many small choices:
 - Which controls to include

Pitfall 1: p-hacking and forking paths

- Garden of forking paths (Gelman & Loken 2013)
- Every analysis involves many small choices:
 - Which controls to include
 - How to code the treatment

Pitfall 1: p-hacking and forking paths

- Garden of forking paths (Gelman & Loken 2013)
- Every analysis involves many small choices:
 - Which controls to include
 - How to code the treatment
 - Which subsample to focus on

Pitfall 1: p-hacking and forking paths

- Garden of forking paths (Gelman & Loken 2013)
- Every analysis involves many small choices:
 - Which controls to include
 - How to code the treatment
 - Which subsample to focus on
 - How to handle outliers

Pitfall 1: p-hacking and forking paths

- Garden of forking paths (Gelman & Loken 2013)
- Every analysis involves many small choices:
 - Which controls to include
 - How to code the treatment
 - Which subsample to focus on
 - How to handle outliers
- If any of those choices were made *after looking at the results*, your p-values are too small

Pitfall 1: p-hacking and forking paths

- Garden of forking paths (Gelman & Loken 2013)
- Every analysis involves many small choices:
 - Which controls to include
 - How to code the treatment
 - Which subsample to focus on
 - How to handle outliers
- If any of those choices were made *after looking at the results*, your p-values are too small
- Mitigations: pre-analysis plans, specification curves, honest reporting of all tried specifications

Pitfall 2: bad controls

- “More controls” is not always better

Pitfall 2: bad controls

- “More controls” is not always better
- **Do not control for:**

Pitfall 2: bad controls

- “More controls” is not always better
- **Do not control for:**
 - Mediators (blocks the effect you want to measure)

Pitfall 2: bad controls

- “More controls” is not always better
- **Do not control for:**
 - Mediators (blocks the effect you want to measure)
 - Colliders (opens a spurious path)

Pitfall 2: bad controls

- “More controls” is not always better
- **Do not control for:**
 - Mediators (blocks the effect you want to measure)
 - Colliders (opens a spurious path)
 - Post-treatment variables (same problem, often both)

Pitfall 2: bad controls

- “More controls” is not always better
- **Do not control for:**
 - Mediators (blocks the effect you want to measure)
 - Colliders (opens a spurious path)
 - Post-treatment variables (same problem, often both)
 - Proxies for the treatment (collinearity)

Pitfall 2: bad controls

- “More controls” is not always better
- **Do not control for:**
 - Mediators (blocks the effect you want to measure)
 - Colliders (opens a spurious path)
 - Post-treatment variables (same problem, often both)
 - Proxies for the treatment (collinearity)
- Every control needs a justification from your DAG

Pitfall 2: bad controls

- “More controls” is not always better
- **Do not control for:**
 - Mediators (blocks the effect you want to measure)
 - Colliders (opens a spurious path)
 - Post-treatment variables (same problem, often both)
 - Proxies for the treatment (collinearity)
- Every control needs a justification from your DAG
- Cinelli, Forney & Pearl (2024) “A crash course in good and bad controls”

Pitfall 3: ignoring dependence

- Standard errors assume observations are independent

Pitfall 3: ignoring dependence

- Standard errors assume observations are independent
- In social science data they usually are not:

Pitfall 3: ignoring dependence

- Standard errors assume observations are independent
- In social science data they usually are not:
 - Observations nested within countries, schools, firms

Pitfall 3: ignoring dependence

- Standard errors assume observations are independent
- In social science data they usually are not:
 - Observations nested within countries, schools, firms
 - Same unit observed over time (panel)

Pitfall 3: ignoring dependence

- Standard errors assume observations are independent
- In social science data they usually are not:
 - Observations nested within countries, schools, firms
 - Same unit observed over time (panel)
 - Spatial proximity

Pitfall 3: ignoring dependence

- Standard errors assume observations are independent
- In social science data they usually are not:
 - Observations nested within countries, schools, firms
 - Same unit observed over time (panel)
 - Spatial proximity
 - Network connections

Pitfall 3: ignoring dependence

- Standard errors assume observations are independent
- In social science data they usually are not:
 - Observations nested within countries, schools, firms
 - Same unit observed over time (panel)
 - Spatial proximity
 - Network connections
- Consequence: default SEs can be *much* too small

Pitfall 3: ignoring dependence

- Standard errors assume observations are independent
- In social science data they usually are not:
 - Observations nested within countries, schools, firms
 - Same unit observed over time (panel)
 - Spatial proximity
 - Network connections
- Consequence: default SEs can be *much* too small
- Fixes: cluster-robust SEs, Conley SEs (spatial), network-adjusted SEs

What's next (not in this class)

- **Multilevel / hierarchical models** (Gelman & Hill)
 - Natural framework when data are nested in groups
- **Bayesian inference**
 - Priors, posteriors, and what MCMC actually does
 - McElreath, *Statistical Rethinking*
- **Machine learning for social science**
 - Prediction-first tools: lasso, random forests, boosting
 - Causal ML (honest trees, double/debiased ML)
- **Text as data**
 - From bag-of-words to LLMs in social science

A checklist for your final essay

- What is the **question** (descriptive, predictive, causal)?

A checklist for your final essay

- What is the **question** (descriptive, predictive, causal)?
- What is the **unit** of analysis and where does **variation** come from?

A checklist for your final essay

- What is the **question** (descriptive, predictive, causal)?
- What is the **unit** of analysis and where does **variation** come from?
- What is the **DAG** behind the relationship?

A checklist for your final essay

- What is the **question** (descriptive, predictive, causal)?
- What is the **unit** of analysis and where does **variation** come from?
- What is the **DAG** behind the relationship?
- What method matches the question and the variation?

A checklist for your final essay

- What is the **question** (descriptive, predictive, causal)?
- What is the **unit** of analysis and where does **variation** come from?
- What is the **DAG** behind the relationship?
- What method matches the question and the variation?
- What are the **identifying assumptions**, and how plausible are they?

A checklist for your final essay

- What is the **question** (descriptive, predictive, causal)?
- What is the **unit** of analysis and where does **variation** come from?
- What is the **DAG** behind the relationship?
- What method matches the question and the variation?
- What are the **identifying assumptions**, and how plausible are they?
- What **alternative explanations** does the design *not* rule out?

A checklist for your final essay

- What is the **question** (descriptive, predictive, causal)?
- What is the **unit** of analysis and where does **variation** come from?
- What is the **DAG** behind the relationship?
- What method matches the question and the variation?
- What are the **identifying assumptions**, and how plausible are they?
- What **alternative explanations** does the design *not* rule out?
- What **robustness** / **placebo** checks can you run?

A checklist for your final essay

- What is the **question** (descriptive, predictive, causal)?
- What is the **unit** of analysis and where does **variation** come from?
- What is the **DAG** behind the relationship?
- What method matches the question and the variation?
- What are the **identifying assumptions**, and how plausible are they?
- What **alternative explanations** does the design *not* rule out?
- What **robustness / placebo** checks can you run?
- How do you **measure** your key concepts, and what's lost in that translation?

Good luck with the essay.

Appendix follows: multilevel models and text as data.

Self-study material, not covered in class.

Roadmap

Appendix: Multilevel models

Appendix: Text as data

Hierarchical data

- Very common structure in social science:

Hierarchical data

- Very common structure in social science:
 - Students within schools within districts

Hierarchical data

- Very common structure in social science:
 - Students within schools within districts
 - Voters within precincts within countries

Hierarchical data

- Very common structure in social science:
 - Students within schools within districts
 - Voters within precincts within countries
 - Respondents within country-years

Hierarchical data

- Very common structure in social science:
 - Students within schools within districts
 - Voters within precincts within countries
 - Respondents within country-years
- Two facts:

Hierarchical data

- Very common structure in social science:
 - Students within schools within districts
 - Voters within precincts within countries
 - Respondents within country-years
- Two facts:
 - Units within the same group are more alike than units in different groups

Hierarchical data

- Very common structure in social science:
 - Students within schools within districts
 - Voters within precincts within countries
 - Respondents within country-years
- Two facts:
 - Units within the same group are more alike than units in different groups
 - Groups vary in how they respond to predictors

Hierarchical data

- Very common structure in social science:
 - Students within schools within districts
 - Voters within precincts within countries
 - Respondents within country-years
- Two facts:
 - Units within the same group are more alike than units in different groups
 - Groups vary in how they respond to predictors
- Ordinary OLS ignores both

Three ways to handle groups

- **Complete pooling:** ignore groups entirely, one intercept
 - Wastes the group-level information
- **No pooling:** separate intercept for each group (fixed effects)
 - Fits each group exactly — but noisily for small groups
- **Partial pooling:** each group gets its own intercept, but shrunk toward a global mean
 - Best of both worlds — this is what multilevel models do

Random intercepts

$$y_{ij} = \alpha_j + \beta x_{ij} + \varepsilon_{ij}, \quad \alpha_j \sim \mathcal{N}(\mu_\alpha, \sigma_\alpha^2)$$

- Each group j has its own intercept α_j
- But the α_j 's are drawn from a common distribution
- Estimate: one global mean + one variance, not 200 separate intercepts
- Versus classic FE: FE *estimates* each α_j ; random intercepts *model* them

Random slopes

$$y_{ij} = \alpha_j + \beta_j x_{ij} + \varepsilon_{ij}$$

- Now *both* intercept and slope vary by group
- Effect of x differs across schools / countries / firms
- Heterogeneity is explicitly modeled, not averaged away
- Useful when you expect the effect itself to depend on context

Fitting multilevel models in R

- Frequentist / REML: `lme4`
 - `lmer(y ~ x + (1 | group), data = d)` — random intercept
 - `lmer(y ~ x + (1 + x | group), data = d)` — random slope
- Bayesian: `brms` or `rstanarm`
 - Same formula syntax, full posterior draws
- Visualization: `sjPlot::plot_model()`, `broom.mixed`

Multilevel vs. fixed effects

	Multilevel	Fixed effects
Goal	Describe variation across groups	Remove group-level confounding
Assumes	Group effects uncorrelated with x	Nothing about that
Small groups	Handles well (shrinkage)	Noisy
Causal setting	Use if assumption plausible	Safer default

- If your goal is descriptive and you have many small groups, multilevel is natural
- If your goal is causal and you worry about group-level confounding, FE is safer

Multilevel: where to go next

- Gelman & Hill (2007) *Data Analysis Using Regression and Multilevel/Hierarchical Models*
- Gelman, Hill & Vehtari (2021) *Regression and Other Stories*, ch. 21–22
- McElreath *Statistical Rethinking*, ch. 13–14
- Bolker's GLMM FAQ:
bbolker.github.io/mixedmodels-misc/glmmFAQ.html

Roadmap

Appendix: Multilevel models

Appendix: Text as data

Why text?

- Political science has always worked with text

Why text?

- Political science has always worked with text
 - Speeches, manifestos, news, laws, tweets, transcripts

Why text?

- Political science has always worked with text
 - Speeches, manifestos, news, laws, tweets, transcripts
- Historically: close reading and hand-coding

Why text?

- Political science has always worked with text
 - Speeches, manifestos, news, laws, tweets, transcripts
- Historically: close reading and hand-coding
- Now: the volume makes automation essential

Why text?

- Political science has always worked with text
 - Speeches, manifestos, news, laws, tweets, transcripts
- Historically: close reading and hand-coding
- Now: the volume makes automation essential
- A quantitative approach: turn documents into features, treat them like any other variable

Preprocessing

- Tokenization: split text into words or sub-words

Preprocessing

- Tokenization: split text into words or sub-words
- Lowercasing

Preprocessing

- Tokenization: split text into words or sub-words
- Lowercasing
- Remove punctuation, numbers, stopwords (“the” , ” and”)

Preprocessing

- Tokenization: split text into words or sub-words
- Lowercasing
- Remove punctuation, numbers, stopwords ("the", "and")
- Stemming / lemmatization: "running", "ran", "runs" → "run"

Preprocessing

- Tokenization: split text into words or sub-words
- Lowercasing
- Remove punctuation, numbers, stopwords (“the”, “and”)
- Stemming / lemmatization: “running”, “ran”, “runs” → “run”
- Every decision changes the downstream analysis

Preprocessing

- Tokenization: split text into words or sub-words
- Lowercasing
- Remove punctuation, numbers, stopwords (“the”, “and”)
- Stemming / lemmatization: “running”, “ran”, “runs” → “run”
- Every decision changes the downstream analysis
- Document these choices the way you would any other transformation

Bag of words

- Simplest representation: count how often each word appears

Bag of words

- Simplest representation: count how often each word appears
- Build a **document-term matrix (DTM)**

Bag of words

- Simplest representation: count how often each word appears
- Build a **document-term matrix (DTM)**
 - Rows: documents

Bag of words

- Simplest representation: count how often each word appears
- Build a **document-term matrix (DTM)**
 - Rows: documents
 - Columns: unique words

Bag of words

- Simplest representation: count how often each word appears
- Build a **document-term matrix (DTM)**
 - Rows: documents
 - Columns: unique words
 - Cells: counts (or TF-IDF weights)

Bag of words

- Simplest representation: count how often each word appears
- Build a **document-term matrix (DTM)**
 - Rows: documents
 - Columns: unique words
 - Cells: counts (or TF-IDF weights)
- Ignores word order and grammar

Bag of words

- Simplest representation: count how often each word appears
- Build a **document-term matrix (DTM)**
 - Rows: documents
 - Columns: unique words
 - Cells: counts (or TF-IDF weights)
- Ignores word order and grammar
- Surprisingly effective for many tasks

Dictionary methods

- Define a list of words associated with a concept

Dictionary methods

- Define a list of words associated with a concept
 - Sentiment: “good”, “great”, “bad”...

Dictionary methods

- Define a list of words associated with a concept
 - Sentiment: “good”, “great”, “bad”...
 - Populism: “elite”, “people”, “corrupt”...

Dictionary methods

- Define a list of words associated with a concept
 - Sentiment: “good”, “great”, “bad”...
 - Populism: “elite”, “people”, “corrupt”...
- Count occurrences per document, maybe normalize

Dictionary methods

- Define a list of words associated with a concept
 - Sentiment: “good”, “great”, “bad”...
 - Populism: “elite”, “people”, “corrupt”...
- Count occurrences per document, maybe normalize
- Pro: simple, transparent, reproducible

Dictionary methods

- Define a list of words associated with a concept
 - Sentiment: “good”, “great”, “bad”...
 - Populism: “elite”, “people”, “corrupt”...
- Count occurrences per document, maybe normalize
- Pro: simple, transparent, reproducible
- Con: brittle to context, sarcasm, domain shift

Dictionary methods

- Define a list of words associated with a concept
 - Sentiment: “good”, “great”, “bad” . . .
 - Populism: “elite”, “people”, “corrupt” . . .
- Count occurrences per document, maybe normalize
- Pro: simple, transparent, reproducible
- Con: brittle to context, sarcasm, domain shift
- Classic in political science — see Laver, Benoit & Garry (2003)
Wordscores; LIWC

Topic models

- Unsupervised: discover themes in a corpus

Topic models

- Unsupervised: discover themes in a corpus
- Latent Dirichlet Allocation (LDA) and descendants

Topic models

- Unsupervised: discover themes in a corpus
- Latent Dirichlet Allocation (LDA) and descendants
 - Each document is a mixture of topics

Topic models

- Unsupervised: discover themes in a corpus
- Latent Dirichlet Allocation (LDA) and descendants
 - Each document is a mixture of topics
 - Each topic is a distribution over words

Topic models

- Unsupervised: discover themes in a corpus
- Latent Dirichlet Allocation (LDA) and descendants
 - Each document is a mixture of topics
 - Each topic is a distribution over words
- Structural Topic Model (STM, Roberts et al.)

Topic models

- Unsupervised: discover themes in a corpus
- Latent Dirichlet Allocation (LDA) and descendants
 - Each document is a mixture of topics
 - Each topic is a distribution over words
- Structural Topic Model (STM, Roberts et al.)
 - Includes document-level covariates (author, year, party)

Topic models

- Unsupervised: discover themes in a corpus
- Latent Dirichlet Allocation (LDA) and descendants
 - Each document is a mixture of topics
 - Each topic is a distribution over words
- Structural Topic Model (STM, Roberts et al.)
 - Includes document-level covariates (author, year, party)
- R packages: `topicmodels`, `stm`

Supervised classification

- You have labels for some documents; predict labels for the rest

Supervised classification

- You have labels for some documents; predict labels for the rest
- Examples:

Supervised classification

- You have labels for some documents; predict labels for the rest
- Examples:
 - Is this speech about foreign policy? (yes / no)

Supervised classification

- You have labels for some documents; predict labels for the rest
- Examples:
 - Is this speech about foreign policy? (yes / no)
 - Is this tweet hate speech? (multi-class)

Supervised classification

- You have labels for some documents; predict labels for the rest
- Examples:
 - Is this speech about foreign policy? (yes / no)
 - Is this tweet hate speech? (multi-class)
- Pipeline: DTM or embeddings → classifier (logistic, SVM, random forest)

Supervised classification

- You have labels for some documents; predict labels for the rest
- Examples:
 - Is this speech about foreign policy? (yes / no)
 - Is this tweet hate speech? (multi-class)
- Pipeline: DTM or embeddings → classifier (logistic, SVM, random forest)
- Evaluation: held-out set, cross-validation, F1/precision/recall

Embeddings and LLMs

- Represent words (or documents) as dense vectors that encode meaning

Embeddings and LLMs

- Represent words (or documents) as dense vectors that encode meaning
- word2vec, GloVe, sentence-BERT...

Embeddings and LLMs

- Represent words (or documents) as dense vectors that encode meaning
- word2vec, GloVe, sentence-BERT...
- Captures word *similarity* (“king” - “man” + “woman” $\approx$ “queen”)

Embeddings and LLMs

- Represent words (or documents) as dense vectors that encode meaning
- word2vec, GloVe, sentence-BERT...
- Captures word *similarity* (“king” - “man” + “woman” $\approx$ “queen”)
- Large language models (GPT, Llama, etc.)

Embeddings and LLMs

- Represent words (or documents) as dense vectors that encode meaning
- word2vec, GloVe, sentence-BERT...
- Captures word *similarity* (“king” - “man” + “woman” $\approx$ “queen”)
- Large language models (GPT, Llama, etc.)
 - Zero-shot classification, extraction, annotation

Embeddings and LLMs

- Represent words (or documents) as dense vectors that encode meaning
- word2vec, GloVe, sentence-BERT...
- Captures word *similarity* (“king” - “man” + “woman” $\approx$ “queen”)
- Large language models (GPT, Llama, etc.)
 - Zero-shot classification, extraction, annotation
 - Replacing hand-coders for many tasks

Embeddings and LLMs

- Represent words (or documents) as dense vectors that encode meaning
- word2vec, GloVe, sentence-BERT...
- Captures word *similarity* (“king” - “man” + “woman” $\approx$ “queen”)
- Large language models (GPT, Llama, etc.)
 - Zero-shot classification, extraction, annotation
 - Replacing hand-coders for many tasks
 - But: black box, bias, replication concerns

Text as data: where to go next

- Grimmer, Roberts & Stewart (2022) *Text as Data*. Princeton UP
 - The standard reference for social scientists
- R packages worth knowing:
 - `quanteda` — full preprocessing & analysis pipeline
 - `stm` — structural topic models
 - `text` — interface to transformer embeddings
- Course: Pablo Barberá's text-as-data materials (online)

Thanks, and good luck.